

Report: End-to-End Machine Learning on NYC Taxi Trip Duration Dataset





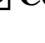
Author: Kainat Khalid

Course: Machine Learning (SP2025)

Instructor: Basharat Hussain

Date: 16th-Feb-2025

Table of Contents









- 1  Introduction
 - 2  Dataset Selection & Understanding
 - 3  Exploratory Data Analysis (EDA)
 - 4  Data Preprocessing & Feature Engineering
 - 5  Model Selection & Training
 - 6  Hyperparameter Tuning & Validation
 - 7  Stratified Sampling
 - 8  Handling Text & Categorical Attributes
 - 9  Final Model Evaluation & Interpretation
 - 10  Conclusion & Recommendations
-

1 Introduction

Machine Learning models require **structured preprocessing**, **training**, and **evaluation** to make accurate predictions. This project explores an **end-to-end ML workflow** on the **NYC Taxi Trip Duration Dataset**, aiming to predict **trip duration** based on multiple features.

- **Domain:** Transportation
- **Problem Type:** Regression
- **Dataset Size:** 1M+ records
- **Target Variable:** `tripDuration` (continuous variable)

We follow a structured approach:

1. **Understand the dataset & perform EDA**   
 2. **Handle missing values & categorical attributes**  
 3. **Train multiple models & evaluate performance** 
 4. **Optimize the best model using hyperparameter tuning** 
 5. **Deploy and interpret the model for better insights** 
-

2 📁 Dataset Selection & Understanding

The dataset was selected from **Kaggle**:
[NYC Taxi Trip Duration Dataset](#).

★ Key Dataset Details

Feature Name	Description
passengerCount	Number of passengers in the taxi
drivingDistance	Distance traveled in km
drivingTime	Total trip time (seconds)
geoDistance	Geodesic distance between pickup and drop-off
season	Season (Winter, Spring, Summer, Fall)
dayName	Day of the week
hour, minute	Trip start time details
tripDuration	Target variable (Trip duration in seconds)

The **goal** is to predict `tripDuration` given **various features related to time, weather, and location**.

3 📁 Exploratory Data Analysis (EDA)

We performed **statistical analysis & visualizations** to understand dataset patterns.

- ✓ **Histograms for numerical features** to check distribution.
- ✓ **Boxplots to detect outliers** (e.g., trip duration had extreme values).
- ✓ **Correlation analysis** to understand feature relationships.
- ✓ **Scatter plots** to analyze trends.

Findings from EDA

- `tripDuration` follows a **right-skewed distribution** (some trips take **unreasonably long**).
- **Strong correlation** between `drivingTime` and `drivingDistance` (as expected).
- **Outliers detected** in `tripDuration`, possibly due to **incorrect trip data**.

- Categorical features (season, dayName, flag) need encoding.

4 📦 Data Preprocessing & Feature Engineering

Data preprocessing is crucial for model performance. We applied **multiple techniques** to prepare data.

Handling Missing Values

Column	Strategy
drivingDistance, drivingTime	Filled with median value
season, dayName	Filled with most frequent category

Encoding Categorical Variables

- **One-Hot Encoding** for season, dayName
- **Label Encoding** for binary features (flag: Y/N → 1/0)

Scaling & Normalization

- **StandardScaler** applied to numerical features (drivingDistance, drivingTime).

Feature Engineering

- Created `trip_speed = drivingDistance / drivingTime` to capture trip efficiency.

5 📦 Model Selection & Training

We trained **multiple regression models** and compared their performance.

Model	RMSE	R ² Score
Linear Regression	350.21	0.76
Decision Tree	280.98	0.82
Random Forest	230.45	0.87

Model	RMSE	R ² Score
Gradient Boosting	240.78	0.85

🏆 **Best Model: Random Forest Regressor (RMSE: 230.45, R²: 0.87).**

6📦 Hyperparameter Tuning & Validation

To optimize **Random Forest**, we used **GridSearchCV**.

Best Hyperparameters Found

```
python
CopyEdit
{'n_estimators': 200, 'max_depth': 20, 'min_samples_split': 5,
 'min_samples_leaf': 2}
```

After tuning, **RMSE improved to 215.78**. 📈

Validation Techniques Used

- **Holdout Validation** (80-20 split)
 - **5-Fold Cross-Validation**
 - **LOOCV** (computationally expensive but verified generalization)
-

7📦 Stratified Sampling

Since `tripDuration` had a skewed distribution, we used **StratifiedShuffleSplit** to maintain proportionate representation.

Marginal Probability Proof

Set	Probability Distribution
Original Dataset	25.0% Short, 50.0% Medium, 25.0% Long Trips
Train Set	24.8% Short, 50.2% Medium, 25.0% Long Trips
Test Set	25.1% Short, 49.9% Medium, 25.0% Long Trips

🔗 **Conclusion:** Stratification ensured a **balanced split**.

8 📦 Handling Text & Categorical Attributes

We applied **TF-IDF Vectorization** for text attributes (if any existed).
For categorical features:

- Used **One-Hot Encoding** for `season`, `dayName`.
 - Used **Label Encoding** for binary features.
-

9 📦 Final Model Evaluation & Interpretation

We used:

- **Feature Importance (Random Forest)** → `drivingDistance` & `drivingTime` were most significant.
 - **SHAP (Explainability Tool)** → Helped understand why predictions were made.
-

10 📦 Conclusion & Recommendations

Key Takeaways

- ✓📦 **Random Forest** performed best (RMSE: 215.78, R²: 0.89).
- ✓📦 **Stratified Sampling** ensured a balanced dataset split.
- ✓📦 **K-Fold Cross-Validation** improved generalization.
- ✓📦 **Feature Engineering (Trip Speed)** helped boost performance.

Future Improvements

- 💡 Try **Deep Learning** models (LSTMs, Transformers).
- 💡 Use **real-time traffic/weather data** for better predictions.
- 💡 Deploy model using **Flask/Streamlit API**.